



CGNS Proposal for Extension 0051

Implicit FaceCenter Indexing for Unstructured Zones with Explicitly Defined Face Sections

Version 1.0

July 10, 2026

CPEX#	0051
Scope	Drop the <code>PointList/PointRange</code> requirement for <code>FaceCenter</code> data on unstructured zones whose face sections (<code>NGON_n</code> or classical 2-D) form one contiguous element-numbering block
Champion	M. Scot Breitenfeld
Organization	The HDF Group
E-mail	brtnfld@hdfgroup.org
GitHub Issue	TBD (implementation gaps tracked as #964 , #965)
Reference Impl.	TBD
Target Release	CGNS v5.0.0 (major release)
Date First Posted	TBD (not yet posted; internal draft)
SIDS Status	proposed
Filemap Status	proposed
MLL Status	proposed

Contents

1 Abstract	4	10
2 Motivation and Rationale	5	11
2.1 Problem Statement	5	12
2.2 Origin of the Restriction	6	13
2.3 Why Explicit Face Sections Change the Picture	7	14

2.4	Precedent: <code>CellCenter</code> Already Works This Way	7	15
2.5	Existing Implementation Gaps	8	16
2.5.1	<code>cg_sol_write</code> rejects plain <code>FaceCenter</code>	8	17
2.5.2	<code>cgnscheck</code> never reads <code>FlowSolution_t</code> 's point set	9	18
3	Detailed Description	10	19
3.1	Design Overview	10	20
3.2	Definition: Complete Face-Section Coverage	10	21
3.3	Revised SIDS Text: <code>FlowSolution_t</code> (§7.7, Note 5)	11	22
3.4	Revised SIDS Text: <code>DiscreteData_t</code> (§12.4)	12	23
3.5	Revised <code>DataSize[]</code> Structure Function	12	24
3.6	MLL Changes	14	25
3.6.1	New internal helper	14	26
3.6.2	Fix: <code>cg_sol_write</code> allow-list	15	27
3.6.3	Fix: <code>cgnscheck</code> point-set-aware sizing	16	28
3.6.4	No change to on-disk representation	16	29
3.7	Scope / Non-Goals	16	30
4	Backward Compatibility	17	31
5	Examples	18	32
5.1	Example 1: Before — Identity <code>PointRange</code> (current, still valid) . . .	18	33
5.2	Example 2: After — Implicit <code>FaceCenter</code> (this CPEX)	19	34
5.3	Example 3: Reading — Dispatching on Coverage	20	35
5.4	Example 4: Rejected — Non-Contiguous Face Sections	20	36
5.5	Example 5: Accepted — <code>NGON_n</code> Coexisting with a Boundary-Patch Section	21	37 38
5.6	Example 6: Accepted — Classical Mesh, Boundary Faces Only . . .	21	39
6	SIDS Impact	22	40
7	File Mapping Impact	22	41
8	Test Plan	22	42
9	Design Decisions	24	43
9.1	Alternative 1: Restrict the Relaxation to <code>NGON_n</code> Coverage Only . .	24	44
9.2	Alternative 2: New Explicit Marker Node (e.g., <code>FacesComplete</code>) . .	24	45
9.3	Alternative 3 (chosen): Structural Inference from Face-Section Coverage	25	46
9.4	Why <code>NGON_n</code> Takes Precedence in Mixed Zones (Tier Rule)	25	47
9.5	On Scope: Why Not Also Amend <code>ZoneSubRegion_t</code> ?	25	48
10	Open Questions	25	49

11 Summary of API Additions	26	50
12 References	27	51

1 Abstract

SIDS Section 7.7 (`FlowSolution_t`), Note 5, states that for unstructured zones `GridLocation` is restricted to `Vertex` or `CellCenter` unless a `PointList` or `PointRange` is also present. The same restriction appears, near word-for-word, in Section 12.4 (`DiscreteData_t`). This forces every `FaceCenter` solution on an unstructured zone to carry an explicit point set — even when the data spans every face the zone defines, in which case the point set is the identity permutation and conveys no information beyond what is already implied by the zone’s own element numbering.

The key observation is that whenever a face exists as an addressable entity in an unstructured zone at all, it exists as an *element* of a face-type `Elements_t` section — `NGON_n` or a classical 2-D type (`TRI_3`, `QUAD_4`, ...) — and therefore already carries a unique global element number (SIDS §7.3). A face that is defined by no section (e.g. an interior face of a purely `HEXA_8` mesh) has no element number, and *cannot* be addressed by an explicit `PointList` today either: the point set requirement never made such faces reachable, it only re-stated numbering the file already contains for the faces that are reachable. “All faces,” for the purposes of face-centered data, has therefore always meant “all faces the zone explicitly defines”: all faces of the grid when `NGON_n` is used (SIDS §7.3 guarantees `NGON_n` specifies them all), and exactly the section-defined — typically boundary — faces otherwise.

This CPEX proposes a precisely scoped relaxation built on that observation: when an unstructured zone has *complete face-section coverage* (defined in Section 3.2 — its face-type sections form a single contiguous block of the zone’s global element numbering), `GridLocation = FaceCenter` data may be written and read *without* a `PointList/PointRange`, indexed directly against that block, exactly as `CellCenter` is indexed against the zone’s `CellDimension`-order elements. The relaxation is strictly additive: files using the existing identity-point-set workaround remain fully conformant. The relaxation covers classical (non-polyhedral) face sections as well as `NGON_n`: the identity point set is exactly as redundant for a boundary-patch `QUAD_4` block as it is for an `NGON_n` block, for the same element-numbering reasons (Section 2.3).

While researching the current behavior, two independent implementation gaps were found that this CPEX’s reference implementation also addresses: `cg_sol_write` cannot create a `FaceCenter`-located solution directly (Section 2.5.1), and `cgnscheck` never reads a `FlowSolution_t`’s `PointList/PointRange` when validating data-array sizes (Section 2.5.2), so today’s identity-point-set files are not actually verifiable by the reference conformance tool.

2 Motivation and Rationale

2.1 Problem Statement

A common CFD data-storage need is a quantity defined on *every* face of an unstructured mesh (e.g., face-normal velocities, surface fluxes on all element boundaries). For a zone built from NFACE_n/NGON_n polyhedral connectivity, the following tree is representative:

```

1 fluid_zone Zone_t [[27 8 0]]
2 +---ZoneType ZoneType_t "Unstructured"
3 +---NGonElements Elements_t [22 0]
4 | +---ElementRange IndexRange_t [ 1 36]
5 | :
6 +---NFaceElements Elements_t [23 0]
7 | +---ElementRange IndexRange_t [37 44]
8 | :
9 +---CellData FlowSolution_t
10 | +---GridLocation GridLocation_t "CellCenter"
11 | +---Pressure DataArray_t /* 8 values -- no PointList/PointRange */
12 +---FaceData FlowSolution_t
13 | +---GridLocation GridLocation_t "FaceCenter"
14 | +---NormalX DataArray_t /* 36 values */
15 | :
16 | +---PointRange IndexRange_t [[1, 36]]
17 |     ~~~~ identity range -- required by Note 5, conveys no new
        information

```

CellData needs no point set: GridLocation = CellCenter alone is sufficient, because the zone's cells are already uniquely and completely numbered by its volume Elements_t sections. FaceData, however, is required by SIDS Note 5 to carry a PointRange, even though [1,36] is exactly the NGonElements section's own ElementRange — information the reader already has.

The same redundancy occurs, for the same reason, on classical (non-polyhedral) meshes. Consider a regular 5×5×5 HEXA_8 zone whose boundary faces are exposed as a QUAD_4 section — the standard layout for attaching BC_t definitions:

```

1 fluid_zone Zone_t [[216 125 0]]
2 +---ZoneType ZoneType_t "Unstructured"
3 +---Hexas Elements_t [17 0] /* HEXA_8 */
4 | +---ElementRange IndexRange_t [ 1 125]
5 | :
6 +---BoundaryQuads Elements_t [7 0] /* QUAD_4 */
7 | +---ElementRange IndexRange_t [126 275]
8 | :
9 +---CellData FlowSolution_t
10 | +---GridLocation GridLocation_t "CellCenter"
11 | +---Pressure DataArray_t /* 125 values -- no PointList/PointRange */
12 +---FaceData FlowSolution_t
13 | +---GridLocation GridLocation_t "FaceCenter"
14 | +---NormalX DataArray_t /* 150 values -- every defined face */
15 | :

```

```

16 +---PointRange IndexRange_t [[126, 275]]
17 ~~~~ identity range again -- the QUAD_4 section's own ElementRange

```

Here the 150-entry solution covers every face the zone *defines*. The zone’s 300 interior faces are not addressable at all — they have no element numbers, so no **PointList** could reference them today either; the point-set requirement has never made them reachable. For the faces that *are* reachable, [126,275] is again exactly the face section’s own **ElementRange**, restated by hand. Two consequences follow in both the polyhedral and the classical case:

1. **Confusion.** Users and tool authors reasonably ask why an identity point set is mandatory for one **GridLocation** and forbidden (well, unnecessary) for another, when both are governed by the same element-numbering machinery.
2. **Detection cost.** Any code path that wants to special-case “this solution spans every face” — to avoid gather/scatter overhead in a parallel read, for instance — must first read the point set and compare it against the face count, on every zone, for every solution. This is a small but repeated and entirely avoidable per-solution check.

2.2 Origin of the Restriction

Note 5 was introduced by CPEX 0031 (“revamped usage of **PointList** and **PointRange**”), incorporated into SIDS v3.1. The SIDS gives its own rationale immediately following the notes list in Section 7.7:

*“For unstructured grids, the value of **GridLocation** alone specifies location and indexing of flow solution data only for vertex and cell-centered data. The reason for this is that element-based grid connectivity provided in the **Elements_t** data structures explicitly indexes only vertices and cells. For data stored at alternate grid locations (e.g., edges), additional connectivity information is needed. This is provided by the optional fields **PointRange** and **PointList**...”*

This is a correct and still-necessary rule for the *general* case: an unstructured zone is only ever *required* to number its vertices and its **CellDimension**-order elements. A zone built purely from **TETRA_4**/**PYRA_5**/**PENTA_6**/**HEXA_8**/**MIXED** sections, with no face-type **Elements_t** section at all, has no face numbering whatsoever — **FaceCenter** is genuinely meaningless there without an explicit point set.

However, the rationale’s premise — that **Elements_t** explicitly indexes “only vertices and cells” — is simply not true of a zone that carries face-type sections. Any **TRI_3**/**QUAD_4** boundary section, and any **NGON_n** section, explicitly indexes faces: that is what an element section *is*. Notably, SIDS v3.1 is also the version in which **NFACE_n** was formalized and the **NGON_n** description was substantially revised (per the same changelog entry). The changes were made in the same release but were not

cross-referenced: Note 5 was written for, and still correctly describes, the zone-with-
no-face-sections case, but it was never revisited to carve out the cases — NGON_n or
classical — where a usable face numbering *does* exist in the file.

2.3 Why Explicit Face Sections Change the Picture

SIDS Section 7.3 (`Elements_t`) already gives every element of every section — faces
included — a unique global number:

*“The elements are indexed with a global numbering system, starting at 1, for all
element sections under a given `Zone_t` data structure. The global numbering
insures that each element, whether it’s a cell, a face, or an edge, is uniquely
identified by its number. They are also listed as a continuous list of element
numbers within any single element section.”*

This guarantee is type-agnostic: a `QUAD_4` boundary section’s faces are exactly as
uniquely and contiguously numbered as an `NGON_n` section’s. The only thing that
varies by element type is *which* faces get numbers. For `NGON_n`, §7.3 additionally
guarantees completeness over the whole grid:

*“Arbitrary polyhedral elements may be defined using the `NGON_n` and `NFACE_n`
element types. **The `NGON_n` element type is used to specify all the faces
in the grid,** and the `NFACE_n` element type is then used to define the polyhedral
elements as a collection of these faces.”*

For classical 2-D sections there is no such completeness statement — they number
only the faces they define, conventionally the boundary. But completeness over the
whole grid was never what the point-set mechanism delivered either: an explicit
`PointList` can only reference faces that have element numbers, so on a classical
mesh it, too, can only ever address the section-defined faces (Section 2.1’s second
example). The set of faces reachable by `FaceCenter` data is *identical* with or without
a point set — it is always “the faces the zone’s sections define.” What the point
set adds is only a restatement of *which contiguous element-ID block* that is, and
whenever the zone’s face sections form a single contiguous block, that restatement
is derivable from the `ElementRanges` the reader already has. This is structurally
identical to how a zone’s volume `Elements_t` sections already give `CellCenter` an
implicit, unambiguous numbering. The `NGON_n` completeness guarantee remains
meaningful, but its role is interpretive, not structural: it tells the reader the block
covers *all* faces of the grid, rather than the boundary subset; the indexing machinery
is the same in both cases.

2.4 Precedent: CellCenter Already Works This Way

SIDS Section 7.7 defines the `DataSet[]` structure function, which is the authoritative
statement of how `FlowSolution_t` array sizes are derived. Quoted verbatim from

the current SIDS source (non-Rind branches; the Rind-present branch follows the same pattern):

```

1  if (PointRange/PointList is present) then
2  {
3    DataSize[] = ListLength[] ;
4  }
5  else if (Rind is absent) then
6  {
7    if (GridLocation = Vertex) or (GridLocation is absent)
8    {
9      DataSize[] = VertexSize[] ;
10   }
11   else if (GridLocation = CellCenter) then
12   {
13     DataSize[] = CellSize[] ;
14   }
15 }

```

The `CellCenter` branch already relies *entirely* on implicit element-based sizing: no `PointList/PointRange` is required there, because `CellSize[]` is already well-defined from the zone's `CellDimension-order Elements_t` sections. This CPEX asks for nothing more than extending that same, already-accepted pattern to `FaceCenter` wherever the zone's face-type sections provide an equally unambiguous numbering — a single contiguous element-ID block (Section 3.2). See Section 3.5 for the precise proposed addition to `DataSize[]`.

2.5 Existing Implementation Gaps

Two related gaps in the reference implementation were found while preparing this proposal. Neither is created by this CPEX, but both must be fixed as part of its reference implementation, since they sit directly in the code paths this CPEX touches.

2.5.1 cg_sol_write rejects plain FaceCenter

`cg_sol_write` (`cgnslib.c`) validates its location argument against an explicit allow-list:

```

1  if (location != CGNS_ENUMV(Vertex) &&
2      location != CGNS_ENUMV(CellCenter) &&
3      location != CGNS_ENUMV(IFaceCenter) &&
4      location != CGNS_ENUMV(JFaceCenter) &&
5      location != CGNS_ENUMV(KFaceCenter)) {
6    cgi_error("Given grid location not supported for FlowSolution_t");
7    return CG_ERROR;
8  }

```

Plain `FaceCenter` — the only location an unstructured zone can use for face data —

is not in this list. This does *not* block the *explicit*-point-set form in Example 5.1: `cg_sol_ptset_write()` already creates a `FaceCenter` solution with an accompanying `PointList/PointRange` in a single call today, because it validates via `cgi_check_location` directly and, internally, works around `cg_sol_write`'s allow-list by calling it with `location = Vertex` and overwriting `sol->location` afterward — itself telling evidence that the allow-list is redundant with `cgi_check_location` rather than an independent rule (see Section 3.6.2). The gap that *does* block today is the *point-set-free* form this CPEX introduces: `cg_sol_ptset_write()` requires `npnts > 0`, so it cannot be used to create a bare `FaceCenter` solution with no point set at all, and `cg_sol_write` is the only entry point for that shape — which is exactly where the allow-list rejects it. Today, the only way to create a point-set-free `FaceCenter` solution is the two-step sequence `cg_sol_write(..., CellCenter, ...)` followed by `cg_gridlocation_write(..., FaceCenter)`, which happens to work only because `cg_gridlocation_write` performs its own, separate validation (`cgi_check_location`) that *does* permit `FaceCenter` for `CellDimension ≥ 3`. This is a latent usability bug independent of Note 5, and it directly affects every user who wants to write a point-set-free `FaceCenter` solution today (e.g. the identity-`PointRange` workaround plus later `cg_gridlocation_write` override, or any future use of this CPEX's implicit form), whether or not this CPEX is adopted. Tracked as [CGNS/CGNS #964](#).

2.5.2 cgnscheck never reads FlowSolution_t's point set

`check_solution()` (`cgnscheck.c`) computes the expected data-array size via `get_data_size()`, which contains:

```

1 if (z->type == CGNS_ENUMV(Unstructured)) {
2     error ("grid location %s not valid for unstructured zone",
3           cg_GridLocationName (location));
4     return 0;
5 }

```

This branch is checked *before* `get_data_size()` ever looks at `GridLocation = FaceCenter`, so it fires for *any* `GridLocation` other than `Vertex/CellCenter` on an unstructured zone — including `FaceCenter` with a fully populated, conformant `PointRange`. `check_solution()` never calls `cg_ptset_info/cg_ptset_read` at all, so today's required identity-point-set workaround (Section 2.1) is not merely *unvalidated* by the reference conformance checker — it is *actively misreported*. `error()` is non-fatal but not silent: it prints a counted `ERROR:` line and increments `cgnscheck`'s error tally for every `FaceCenter` solution on every unstructured zone in the file, today, for files that are fully conformant under the *current* SIDS. This is a stronger motivation than mere non-validation: it is a false-positive in the reference tool that this CPEX's fix also happens to resolve. This CPEX's reference implementation fixes `get_data_size()` to honor `ListLength[]` from an explicit point set (already

required by the current SIDS) in addition to the new implicit face-section sizing this CPEX adds. Tracked as [CGNS/CGNS #965](#).

3 Detailed Description

3.1 Design Overview

The proposal has four parts:

1. A formal SIDS definition of *complete face-section coverage* (Section 3.2), built entirely from existing SIDS guarantees — no new node types or structural elements are introduced.
2. A redline amendment to SIDS Note 5 in `FlowSolution_t` (§7.7) and the matching clause in `DiscreteData_t` (§12.4), and a corresponding addition to the `DataSet[]` structure function (Section 3.5).
3. MLL changes (Section 3.6) to compute and validate the implicit face count, expose it to readers via a new public accessor (Section 3.6.1), and fix the two unrelated implementation gaps from Section 2.5 (as standalone, independently backportable patches — see Section 11).
4. A test plan (Section 8) covering the coverage predicate, the new sizing rule, and the write-ordering constraint this CPEX introduces (Section 4).

3.2 Definition: Complete Face-Section Coverage

Call an `Elements_t` section of a `CellDimension-3` unstructured zone a **face section** if its element type is `NGON_n` or a classical 2-D type (`TRI_3`, `TRI_6`, `QUAD_4`, `QUAD_8`, `QUAD_9`, or a `MIXED` section containing only such types). The zone’s **face-indexing sections** are then determined by a two-tier precedence rule:

- T1. If Z contains one or more `NGON_n` sections, the face-indexing sections are the `NGON_n` sections *only*. SIDS §7.3 states that `NGON_n` “is used to specify all the faces in the grid,” so when present it is authoritative; classical face sections coexisting with it (typically boundary-patch duplicates kept for `BC_t`/tool compatibility) are geometric restatements of faces the `NGON_n` block already numbers, and including them would double-count those faces.
- T2. Otherwise, the face-indexing sections are all of Z ’s classical face sections.

An unstructured zone Z has **complete face-section coverage** if and only if:

1. Z has at least one face-indexing section; and
2. the union of the `ElementRanges` of the face-indexing sections forms a single contiguous block $[E_{\text{start}}, E_{\text{end}}]$ of the zone’s global element numbering (trivially true

for one section; for multiple sections, their ranges must partition $[E_{\text{start}}, E_{\text{end}}]$ with no gaps, and no element index inside that block may belong to any section that is not one of the face-indexing sections).

When these conditions hold, the face count is $N_{\text{face}} = E_{\text{end}} - E_{\text{start}} + 1$, and the **implicit FaceCenter indexing rule** is: the j -th entry (1-based) of the `DataArray_t` maps to global element index

$$E_{\text{start}} + j - 1, \quad 1 \leq j \leq N_{\text{face}}.$$

Three points about what this rule does and does not mean:

- **“Complete” means complete over the defined faces.** Under tier T1 the block covers every face of the grid, by NGON_n’s own SIDS guarantee. Under tier T2 it covers every face the zone *explicitly defines* — conventionally the boundary faces. Interior faces of a classical mesh have no element numbers and are equally unreachable by an explicit `PointList` today (Section 2.3); this CPEX neither shrinks nor grows the set of addressable faces, it only removes the redundant restatement of their numbering.
- **The rule refers only to the block.** Sections outside $[E_{\text{start}}, E_{\text{end}}]$ — volume sections, and (under T1) classical boundary-patch sections — do not disqualify Z and do not change what the rule means; they simply are not part of the indexing block. Condition 2 alone guarantees the block is unambiguous regardless of what else exists in the zone.
- **Ambiguity is structurally excluded.** A reader never guesses which section a point-set-free `FaceCenter` array belongs to: the two-tier rule plus the contiguity condition yield either exactly one well-defined block or a determination that the zone does not qualify (in which case an explicit point set remains required, exactly as today). In particular, a zone with several same-typed boundary patches qualifies only if their ranges are mutually contiguous — the array then spans all of them in global element order — and never maps to any single patch by a size coincidence.

This CPEX does not extend to `EdgeCenter`. On `CellDimension-2` zones `FaceCenter` itself is not a valid location (SIDS §7.7’s location table), and `EdgeCenter` data on 2-D or 3-D zones raises the analogous question for `BAR_2`-type sections; the machinery defined here would transfer directly, but that extension is left to a future CPEX rather than folded into this one (Section 3.7).

3.3 Revised SIDS Text: `FlowSolution_t` (§7.7, Note 5)

Current text:

*For unstructured zones GridLocation options are limited to Vertex
or CellCenter, unless one of PointList or PointRange is present.*

Proposed text:

For unstructured zones GridLocation options are limited to Vertex
or CellCenter, unless one of PointList or PointRange is present,
or unless GridLocation is FaceCenter and the zone has complete
face-section coverage, in which case indexing follows the zone's
face-section element numbering (§7.3) directly.

3.4 Revised SIDS Text: DiscreteData_t (§12.4)

The matching clause in DiscreteData_t is amended the same way, for consistency
— both notes were introduced together by CPEX 0031 and have always read the
same way, modulo word order (“PointRange or PointList” here vs. “PointList or
PointRange” in §7.7):

Current text:

*For unstructured zones GridLocation options are limited to Vertex
or CellCenter, unless one of PointRange or PointList is present.*

Proposed text:

For unstructured zones GridLocation options are limited to Vertex
or CellCenter, unless one of PointRange or PointList is present,
or unless GridLocation is FaceCenter and the zone has complete
face-section coverage, in which case indexing follows the zone's
face-section element numbering (§7.3) directly.

ZoneSubRegion_t and BC_t, the other two structures touched by CPEX 0031, are
deliberately *not* amended here. ZoneSubRegion_t always requires an explicit region
descriptor (point set, BCRegionName, or GridConnectivityRegionName) regardless
of GridLocation, so Note 5 there is not a burden in the same way. BC_t does not
carry this clause at all (a BC patch is definitionally a *subset*, so an explicit point
set is appropriate there even for a complete-coverage zone). See Section 9.5 for
discussion of whether a future CPEX should revisit ZoneSubRegion_t as well.

3.5 Revised DataSize[] Structure Function

Section 7.7's DataSize[] definition gains one new FaceCenter case inside *each* of
the two existing Rind branches, in direct parallel with the existing Vertex and
CellCenter cases. The explicit PointRange/PointList branch remains first and
still takes precedence, so existing identity-point-set files are unaffected; and because

the new cases sit *inside* the **Rind** branches (rather than ahead of them), a **FaceCenter** solution carrying rind elements is still sized with its rind contribution, exactly as **CellCenter** is. The redline below follows the existing SIDS pseudocode style (**then** keywords, braces on their own lines) exactly as it appears in the current SIDS source, so the amendment can be applied to the standard’s text verbatim; the two new cases are marked:

```

1  if (PointRange/PointList is present) then
2  {
3    DataSize[] = ListLength[] ;
4  }
5  else if (Rind is absent) then
6  {
7    if (GridLocation = Vertex) or (GridLocation is absent)
8    {
9      DataSize[] = VertexSize[] ;
10   }
11   else if (GridLocation = CellCenter) then
12   {
13     DataSize[] = CellSize[] ;
14   }
15   else if (GridLocation = FaceCenter) and /* new */
16         (zone has complete face-section coverage) then
17   {
18     DataSize[] = FaceSize[] ;
19   }
20 }
21 else if (Rind is present) then
22 {
23   if (GridLocation = Vertex) or (GridLocation is absent) then
24   {
25     DataSize[] = VertexSize[] + [a + b,...] ;
26   }
27   else if (GridLocation = CellCenter)
28   {
29     DataSize[] = CellSize[] + [a + b,...] ;
30   }
31   else if (GridLocation = FaceCenter) and /* new */
32         (zone has complete face-section coverage) then
33   {
34     DataSize[] = FaceSize[] + [a + b,...] ;
35   }
36 }

```

Here **FaceSize[]** is defined as $E_{\text{end}} - E_{\text{start}} + 1$ for the zone’s contiguous face-indexing block (Section 3.2), and **RindPlanes** = [a,b,...] as in the existing definition. The name **FaceSize[]**, rather than **NFace[]**, is chosen to avoid confusion with **NFACE_n** cell counts.

Note that §12.4 requires no matching **DataSize[]** amendment: the SIDS defines **DiscreteData_t**’s **DataSize[]** by reference — “It is identical to the function **DataSize[]** defined for **FlowSolution_t**” — so the change above propagates to **DiscreteData_t** automatically. Only the §12.4 *Note* (Section 3.4) needs its own

redline.

343

3.6 MLL Changes

344

3.6.1 New internal helper

345

```

1  /* cgns_internals.c */
2  int cgi_zone_face_range(cgns_zone *zone,
3                          cgsize_t *start, cgsize_t *end);
4      /* Returns CG_OK and the contiguous [start,end] face-indexing
5       * block if the zone has complete face-section coverage per the
6       * definition in this document's Detailed Description section,
7       * applying the two-tier rule: NGON_n sections define the block
8       * when any are present; otherwise all classical 2-D face
9       * sections do. Returns CG_NODE_NOT_FOUND if the zone has no
10      * face-type sections at all; returns CG_ERROR (with cgi_error()
11      * detail) if face-type sections exist but the applicable tier's
12      * union is not one contiguous block. The cgi_error() message
13      * names the zone and, for the gap case, the interrupting section
14      * (e.g. "zone 'fluid_zone': face sections [1,20] and [31,50]
15      * are not contiguous (elements 21-30 belong to section 'Hexas')",
16      * so the failure is diagnosable without re-deriving element
17      * ranges by hand. Rejects (CG_ERROR) any ElementRange with
18      * start < 1 or end < start, and any face count that would
19      * overflow cgsize_t arithmetic in the caller's size computation,
20      * before returning -- both ranges are read from the file and
21      * must not be trusted blindly. */

```

This helper is used by `cg_sol_write/cg_field_write` (and their `cg_discrete_*` counterparts) whenever `GridLocation = FaceCenter` is requested on an unstructured zone with no accompanying point set, to (a) confirm the relaxation applies and (b) compute the expected `DataArray_t` size for validation on write, exactly as `CellSize[]` is already used for `CellCenter`. Coverage is a property of the zone's `Elements_t` section layout, already resident in memory as `zone->section[]`; the reference implementation computes it once per zone (caching the result on `cgns_zone`) rather than re-scanning the section list on every `cg_field_write` call to a multi-field `FaceCenter` solution, and invalidates the cached result on `cg_section_write/cg_section_partial_write`.

A thin public wrapper is also added, so that reader-side code (as in Example 5.3) does not have to reimplement the coverage predicate independently for every downstream tool:

```

1  /* cgnslib.c; public API */
2  int cg_zone_face_range(int fn, int B, int Z,
3                        cgsize_t *start, cgsize_t *end);
4      /* Thin wrapper over cgi_zone_face_range() for read-side
5       * consumers. Same return-value contract. */

```

3.6.2 Fix: cg_sol_write allow-list

359

Rather than add FaceCenter as one more literal in `cg_sol_write`'s local allow-list, the reference implementation *removes* that allow-list and delegates to `cgi_check_location()` — the same dimension- and zone-type-aware check that `cg_gridlocation_write` and `cg_sol_ptset_write` already trust, and the reason those two entry points already accept FaceCenter correctly today while `cg_sol_write` does not. `cg_sol_write`'s own literal allow-list and its separate, hand-written [IJK]FaceCenter/Structured-only check are two independent copies of a rule `cgi_check_location` already enforces; keeping both is exactly the kind of duplication that let them drift apart and produced this bug in the first place. Because `cgi_check_location` needs the zone's `CellDimension` and `ZoneType_t`, the check must move to *after* the zone lookup (`cg_sol_write` currently validates `location` before `cgi_get_file/cgi_get_zone` are called):

371

```

1  /* was: a literal allow-list checked before the zone lookup, plus a
2   * separate structured-only [IJK]FaceCenter check further down. Both
3   * are replaced by one delegated call, placed after the zone lookup: */
4  zone = cgi_get_zone(cg, B, Z);
5  if (zone == 0) return CG_ERROR;
6  if (cgi_check_location(cg->base[B-1].cell_dim, zone->type, location))
7      return CG_ERROR; /* cgi_check_location has already set cgi_error()
                        */

```

This is strictly additive in effect: every `location` value previously accepted by the old allow-list is still accepted (each is one of `cgi_check_location`'s `CG_OK` cases), plain FaceCenter is now accepted for `CellDimension` ≥ 3 regardless of zone type (matching what `cg_gridlocation_write` already permits today), and [IJK]FaceCenter on a non-Structured zone is still rejected, with the same error message, by `cgi_check_location` itself. This fix is independent of the rest of this CPEX (it does not by itself relax Note 5) but is required in order for Example 5.2's single-call write to work; without it, applications must still fall back to the two-step `cg_sol_write` + `cg_gridlocation_write` sequence for the point-set-free form. Being a narrowly-scoped, independently testable bug fix with no dependency on the face-section-coverage logic, it can land as its own patch against the current release line without waiting on the rest of this CPEX; see the note in Section 11.

383

3.6.3 Fix: cgnscheck point-set-aware sizing

`get_data_size()` (Section 2.5.2) is updated to: (1) read a `PointList/PointRange` under `FlowSolution_t/DiscreteData_t`, if present, and use its `ListLength` as the expected size (this branch is required by the *current* SIDS and was simply never implemented); and (2) when no point set is present and `GridLocation = FaceCenter` on an unstructured zone, call `cgi_zone_face_range()` and use the returned count. Only when neither applies does `get_data_size()` retain its current “grid location not valid for unstructured zone” error.

3.6.4 No change to on-disk representation

No new node type, attribute, or data layout is introduced. A conformant file under this CPEX is a `FlowSolution_t` (or `DiscreteData_t`) node with `GridLocation = FaceCenter` and *no* `PointList/PointRange` child, on a zone with complete face-section coverage. Because this is a strict superset of what Note 5 currently forbids, this CPEX changes only the *normative interpretation* of an existing, legal-today node shape — it does not add anything a pre-CPEX HDF5/ADF reader would fail to open at the byte level.

3.7 Scope / Non-Goals

- **Not addressed: EdgeCenter.** The machinery in Section 3.2 would transfer directly to `EdgeCenter` data indexed against `BAR`-type sections — the natural objection (“no completeness guarantee for edges”) does not apply, under the same “complete over the defined elements” reading this CPEX uses for faces — but the extension multiplies the location/dimension test matrix and is deliberately left to a future CPEX rather than folded into this one.
- **Not addressed:** the structured-zone `GridLocation` gap raised during the same pre-submission review: 2-D structured zones have no directional edge-center locations (`IEdgeCenter/JEdgeCenter`) analogous to 3-D’s `IFaceCenter/JFaceCenter/KFaceCenter` — `GridLocation_t` today offers only the undifferentiated `EdgeCenter`. This is a real asymmetry in the enumeration, but it concerns structured zones and a new enumeration value, orthogonal to this CPEX’s unstructured-zone/no-new-enumerations scope; it is recorded here as a candidate companion CPEX (Section 10).
- **Not addressed: ZoneSubRegion_t, BC_t.** See Section 9.5.
- **Not required:** no change to `NFACE_n`’s own definition, or to how `CellCenter` is already handled; both are unaffected and serve as the existing precedent this CPEX generalizes.

4 Backward Compatibility

419

- **Existing files:** Fully compatible. Files using the current identity-point-set
workaround remain valid; the explicit `PointRange/PointList` branch of the
revised `DataSize[]` function still takes precedence over the new implicit branch,
so nothing about their interpretation changes.
- **Existing readers (pre-CPEX-0051).** A pre-CPEX reader that encounters
a `FaceCenter FlowSolution_t` with no point set on a qualifying zone will,
under the *old* SIDS text, consider the file non-conformant (Note 5 as currently
written requires the point set), and a naive reader may simply misinterpret
the array size instead of erroring cleanly. This is the one real compatibility
wrinkle, and it deserves emphasis: writers that adopt this CPEX are pro-
ducing files that are conformant under the *new* SIDS text but not the old
one, so **a writer must not omit the point set unless it knows its
downstream tool ecosystem supports this CPEX**. CGNS versions its
SIDS compliance via `CGNSLibraryVersion_t` (and, if CPEX 0049 is adopted,
`CGNSMinRequiredVersion_t`). If CPEX 0049 is adopted, this CPEX **must**
be registered in its Feature-Version Registry (CPEX 0049, “Feature-Version
Registry” section) so that AUTO write mode raises a file’s minimum required
version whenever the implicit `FaceCenter` form is used, and so that version-
aware readers reject or warn appropriately rather than silently misreading array
sizes. If CPEX 0049 is not adopted, this CPEX should instead mandate a plain
`CGNSLibraryVersion_t` bump on write, as the minimum viable version-gating
mechanism.

Why a major release. This CPEX is targeted at CGNS **v5.0.0**, the next
major release, rather than a point release on the v4.x line. This is not merely
a labeling choice: under CPEX 0049’s revised `cg_open` logic, a library whose
major version is older than the file’s major version is rejected immediately, before
any feature-compatibility check runs (CPEX 0049, “Modified `cg_open` Behavior,”
major-version safety guard). If this CPEX lands in v5.0.0 alongside CPEX 0049,
a v4.x reader opening a v5.0-written file is hard-rejected at that guard and never
reaches the point where it would have to interpret a point-set-free `FaceCenter`
solution at all, eliminating the “silently misreads the data” failure mode entirely
for that population of readers. This is a stronger guarantee than a minor-version
bump could offer under the same mechanism, since CPEX 0049 only hard-rejects
on a major-version mismatch; a same-major, higher-minor file is instead routed
through the feature-aware check, which depends on this CPEX’s Feature-Version
Registry entry being present and correctly detected.
- **Existing writers:** Unaffected. Codes that do not use the new implicit form
continue to write the explicit point set exactly as before; nothing is deprecated or
removed. *Today*, this ordering independence is unconditional: `cg_field_write`

branches on whether a point set is already attached to the solution at the moment *it* is called (`sol->ptset == NULL`), and writing a `FaceCenter` field before any point set exists already fails on an unstructured zone (no implicit sizing rule exists yet), so no existing writer can depend on a field-before-point-set ordering. This CPEX changes that: once implicit sizing exists, a solution written field-first would size against the face-section block count, and a point set added *afterward* would silently reinterpret an already-on-disk array written to a different, incompatible size. The reference implementation must therefore forbid attaching a `PointList/PointRange` to a `FlowSolution_t/ DiscreteData_t` node that already has one or more `DataArray_t` children written under implicit sizing — returning `CG_ERROR` from the point-set write call in that case — rather than silently permitting the reinterpretation. This is a new ordering constraint introduced by this CPEX, not a preexisting one, and should be covered by its own test case (Section 8).

- **ABI:** No structure layout changes are required beyond the addition of the internal `cgi_zone_face_range()` helper and its public wrapper `cg_zone_face_range()` (Section 3.6.1). The public API changes are limited to widening `cg_sol_write`'s accepted `location` values (already a bug fix independent of this CPEX) and adding the new read-side wrapper, and are purely additive in effect (previously-rejected input is now accepted; no previously-accepted input is rejected; the wrapper is an optional convenience with no existing caller). **No new public error codes are introduced.** A write requesting implicit `FaceCenter` sizing on a zone that lacks complete face-section coverage returns the existing `CG_ERROR` (with a `cgi_error()` message identifying the missing coverage), exactly as an out-of-range `GridLocation` does today; `cg_zone_face_range()` returns the same `CG_OK/CG_NODE_NOT_FOUND/CG_ERROR` contract as every other existing `cg_*_info`-style read function, so no new return-code semantics are introduced at the public boundary either.
- **Parallel I/O (`cgp_*`):** Unaffected. The relaxation changes only whether a point-set child node must exist; it does not change the collective/independent write model for `DataArray_t` contents.

5 Examples

5.1 Example 1: Before — Identity `PointRange` (current, still valid)

```

1  int S, F;
2  cg_sol_write(fn, B, Z, "CellData", CGNS_ENUMV(CellCenter), &S);
3  cg_field_write(fn, B, Z, S, CGNS_ENUMV(RealDouble), "Pressure",
4                pressure, &F);
5
6  /* FaceCenter WITH an explicit point set already works today in one
7   * call via cg_sol_ptset_write -- the gap this CPEX's reference
8   * implementation fixes (see Existing Implementation Gaps, above) is
9   * specifically the point-set-FREE form below, not this one. The
10  * identity PointRange must still be kept in sync with NGonElements'
11  * ElementRange by hand. */
12  cgsize_t range[2] = {1, 36};
13  cg_sol_ptset_write(fn, B, Z, "FaceData", CGNS_ENUMV(FaceCenter),
14                    CGNS_ENUMV(PointRange), 2, range, &S);
15  cg_field_write(fn, B, Z, S, CGNS_ENUMV(RealDouble), "NormalX",
16                normal_x, &F);

```

5.2 Example 2: After — Implicit FaceCenter (this CPEX)

492

```

1  int S, F;
2
3  /* FaceCenter now accepted directly by cg_sol_write (fix described in
4   * this CPEX's MLL Changes section); no PointRange needed -- the
5   * zone's NGonElements section already defines the face numbering. */
6  cg_sol_write(fn, B, Z, "FaceData", CGNS_ENUMV(FaceCenter), &S);
7  cg_field_write(fn, B, Z, S, CGNS_ENUMV(RealDouble), "NormalX",
8                normal_x, &F); /* size validated against the zone's
9                               * face-section block count */

```

Resulting tree (compare to the first example in Section 2.1):

493

```

1  fluid_zone Zone_t [[27 8 0]]
2  +---ZoneType ZoneType_t "Unstructured"
3  +---NGonElements Elements_t [22 0]
4  |   +---ElementRange IndexRange_t [1 36]
5  +---NFaceElements Elements_t [23 0]
6  |   +---ElementRange IndexRange_t [37 44]
7  +---CellData FlowSolution_t
8  |   +---GridLocation GridLocation_t "CellCenter"
9  |   +---Pressure DataArray_t
10 +---FaceData FlowSolution_t
11 |   +---GridLocation GridLocation_t "FaceCenter"
12 |   +---NormalX DataArray_t /* 36 values; indexed 1..36 against
13                             NGonElements, no PointRange node */

```

The MLL's Fortran bindings (`cg_sol_write_f`, `cg_field_write_f`) forward the same location argument through the same validation as their C counterparts, so both the `cg_sol_write` allow-list fix and this CPEX's implicit-sizing behavior are directly visible to Fortran callers, who make up a substantial share of CGNS writers. The Fortran mirror of Example 5.2:

494

495

496

497

498

```

1 integer :: S, F, ier
2
3 ! FaceCenter now accepted directly by cg_sol_write_f; no PointRange
4 ! needed -- the zone's NGonElements section defines the numbering.
5 call cg_sol_write_f(fn, B, Z, 'FaceData', FaceCenter, S, ier)
6 if (ier .ne. CG_OK) call cg_error_exit_f()
7 call cg_field_write_f(fn, B, Z, S, RealDouble, 'NormalX', &
8     normal_x, F, ier)
9 if (ier .ne. CG_OK) call cg_error_exit_f()

```

5.3 Example 3: Reading — Dispatching on Coverage

499

```

1 char name[33];
2 CGNS_ENUMT(GridLocation_t) loc;
3 CGNS_ENUMT(PointSetType_t) ptype;
4 csize_t npnts, fstart, fend;
5
6 if (cg_sol_info(fn, B, Z, S, name, &loc)) return handle_error();
7
8 /* cg_sol_ptset_info addresses solution S directly and reports "no
9 * point set" as CG_OK with ptype == PointSetTypeNull -- unlike the
10 * cg_goto-positioned cg_ptset_info, it needs no prior cg_goto call,
11 * and a real read failure is distinguishable from "point set absent". */
12 if (cg_sol_ptset_info(fn, B, Z, S, &ptype, &npnts)) return handle_error();
13 int has_ptset = (ptype != CGNS_ENUMV(PointSetTypeNull));
14
15 if (loc == CGNS_ENUMV(FaceCenter) && !has_ptset) {
16     /* Node shape alone (FaceCenter, no point set) is necessary but
17     * NOT sufficient: a pre-CPEX-0051 or non-conformant file can
18     * present this same shape without complete face-section
19     * coverage. Verify before trusting it -- this CPEX makes the
20     * check cheap, not optional, which is why cg_zone_face_range
21     * is public. */
22     if (cg_zone_face_range(fn, B, Z, &fstart, &fend) != CG_OK) {
23         /* Claims the implicit form but coverage doesn't hold:
24         * non-conformant file. Fail safely rather than guessing a
25         * size -- this is the "silently misreads the data" failure
26         * mode Section 4 warns against. */
27         return handle_nonconformant_file(fn, B, Z, S);
28     }
29     /* Spans every face the zone defines: all faces of the grid on
30     * an NGON_n zone (tier T1), the section-defined (typically
31     * boundary) faces on a classical zone (tier T2). */
32     process_all_defined_faces(fn, B, Z, S, fend - fstart + 1);
33 } else if (loc == CGNS_ENUMV(FaceCenter) && has_ptset) {
34     /* Explicit form: a genuine subset, or a pre-CPEX-0051 file. */
35     process_face_subset(fn, B, Z, S, ptype, npnts);
36 }

```

5.4 Example 4: Rejected — Non-Contiguous Face Sections

500

```

1  /* Case A (polyhedral): TWO NGON_n sections separated by an
2  * unrelated section, e.g. ElementRange [1,20] (NGON_n), [21,30]
3  * (some other type), [31,50] (NGON_n). Condition 2 of the coverage
4  * definition is violated for tier T1: the NGON_n ranges do not form
5  * a single contiguous block, so there is no well-defined
6  * [E_start, E_end] to index against.
7  *
8  * Case B (classical): QUAD_4 patches interleaved with volume
9  * sections, e.g. HEXA_8 [1,100], QUAD_4 [101,150], HEXA_8
10 * [151,200], QUAD_4 [201,250]. Tier T2 applies (no NGON_n), and
11 * the union of the two QUAD_4 ranges is not contiguous. In both
12 * cases the explicit point set is still required -- notably, this
13 * is what structurally prevents a same-sized-patch ambiguity: the
14 * implicit form never attaches to one patch among several by a
15 * size coincidence. */
16 cg_sol_write(fn, B, Z, "FaceData", CGNS_ENUMV(FaceCenter), &S);
17 /* cg_field_write size-validation falls back to requiring an
18 * explicit PointList/PointRange, and cgnscheck flags its absence. */

```

5.5 Example 5: Accepted — NGON_n Coexisting with a Boundary-Patch Section

```

1  /* Zone has an NGON_n section [1,36] covering every face, PLUS a
2  * separate QUAD_4 section [37,44] duplicating the boundary faces
3  * for a BC_t patch -- a common pattern for tool compatibility.
4  * This qualifies under tier T1: NGON_n is present, so it alone
5  * defines the face-indexing block; the QUAD_4 duplicates are
6  * excluded, which is what prevents the boundary faces from being
7  * counted twice (N stays 36, not 44). The QUAD_4 range must merely
8  * lie outside the NGON_n block, which it does. */
9  cg_sol_write(fn, B, Z, "FaceData", CGNS_ENUMV(FaceCenter), &S);
10 cg_field_write(fn, B, Z, S, CGNS_ENUMV(RealDouble), "NormalX",
11               normal_x, &F); /* 36 values, indexed against [1,36] */

```

5.6 Example 6: Accepted — Classical Mesh, Boundary Faces Only

The $5 \times 5 \times 5$ HEXA_8/QUAD_4 zone from Section 2.1: no NGON_n anywhere, one classical face section [126,275].

```

1  /* Tier T2: no NGON_n sections, so the face-indexing block is the
2  * union of the classical face sections -- here the single QUAD_4
3  * section [126,275], trivially contiguous. N_face = 150. */
4  cg_sol_write(fn, B, Z, "FaceData", CGNS_ENUMV(FaceCenter), &S);
5  cg_field_write(fn, B, Z, S, CGNS_ENUMV(RealDouble), "NormalX",
6               normal_x, &F); /* 150 values, indexed against
7                               * [126,275] -- every DEFINED face.
8                               * The 300 interior faces have no
9                               * element numbers and remain exactly
10                              * as unaddressable as they are today
11                              * with an explicit point set. */

```

Multiple boundary patches also qualify when contiguous: three QUAD_4 sec-

tions [126,175], [176,225], [226,275] (e.g. INFLOW/OUTFLOW/WALL) partition [126,275] with no gaps, so the zone qualifies with the same $N_{\text{face}} = 150$; the implicit array spans all three patches in global element order. Per-patch data still uses an explicit point set, exactly as today.

6 SIDS Impact

- §7.7 FlowSolution_t, Note 5: amended (Section 3.3).
- §7.7 FlowSolution_t, DataSize[] function: amended (Section 3.5).
- §12.4 DiscreteData_t, equivalent Note: amended (Section 3.4). No DataSize[] amendment is needed in §12.4 itself: the SIDS defines DiscreteData_t's DataSize[] by reference ("It is identical to the function DataSize[] defined for FlowSolution_t"), so the §7.7 change propagates automatically (Section 3.5).
- §7.3 Elements_t: unchanged. This CPEX relies entirely on the existing global-numbering guarantee (all element types) and, for the tier rule, the "NGON_n specifies all faces" statement already there (Section 2.3); no new text is needed in §7.3 itself.
- No new SIDS node types, enumerations, or structural fields are introduced anywhere.

7 File Mapping Impact

None. This CPEX changes only which combinations of existing nodes are considered conformant; it introduces no new HDF5/ADF on-disk structures, attributes, or data types. The ADFH and ADF file-mapping documents require no amendment.

8 Test Plan

The reference implementation's test suite must cover at least the following fixtures. Items 1–7 exercise the coverage predicate (both tiers) and the new sizing rule directly; items 8–9 pin down behavior this CPEX makes newly reachable and that would otherwise be left implicit; item 10 guards the one existing entry point this CPEX widens.

1. Single NGON_n section, implicit form (Example 5.2): round-trip write, read, and cgnscheck clean (zero errors) in both C and Fortran (Fortran mirror of Example 5.2).

-
2. Multiple contiguous NGON_n sections partitioning one [E_start, E_end] block: 538
accepted, sized correctly. 539
 3. Classical-only zone, tier T2 (Example 5.6): single QUAD_4 boundary section 540
on a HEXA_8 mesh — round-trip write, read, cgnscheck clean, N_{face} equal to 541
the section's ElementSize. Repeat with three contiguous same-typed patches 542
partitioning the block, and with a mixed TRI_3+QUAD_4 contiguous pair. *This* 543
is entirely new behavior with no installed base to regress against, so these fixtures 544
carry the highest regression value in the suite alongside item 4. 545
 4. Mixed zone, tier T1 precedence (Example 5.5): NGON_n block plus classical 546
boundary-patch duplicates outside it — accepted, and N_{face} must equal the 547
NGON_n block size alone (36, not 44, in the example's numbers), pinning down 548
that classical duplicates are excluded and boundary faces are not double-counted. 549
 5. Gap cases (Example 5.4, both Case A and Case B): write of a solution requesting 550
implicit sizing is rejected with CG_ERROR; a file constructed to have this shape 551
externally is flagged by cgnscheck and by cg_zone_face_range() returning 552
CG_ERROR on read. Case B doubles as the same-sized-patch ambiguity guard: 553
two disjoint equal-length QUAD_4 patches must never be silently resolved to 554
either patch. 555
 6. Legacy identity-PointRange file (Example 5.1): still passes, unchanged size, 556
confirming the explicit-point-set branch of DataSize[] retains precedence over 557
the new implicit branch. Repeat for the classical-mesh identity file of Section 2.1 558
(PointRange [126,275]). 559
 7. Zone with no face-type sections at all (pure HEXA_8/MIXED-volume): cg_zone_- 560
face_range() returns CG_NODE_NOT_FOUND; implicit FaceCenter write is re- 561
jected; the explicit-point-set path remains the only one, exactly as today. 562
 8. Point-set-written-after-implicit-fields ordering: attempt to attach a 563
PointRange/PointList to a FlowSolution_t that already holds fields 564
written under implicit face-section sizing; must return CG_ERROR per the rule in 565
Section 4. This behavior is otherwise undefined by the prose alone; the test is 566
what pins it down. 567
 9. Malformed ElementRange inputs to cgi_zone_face_range(): start < 1, 568
end < start, and (on a 32-bit cgsized_t build) a range whose face count 569
overflows subsequent size arithmetic — all three must return CG_ERROR before 570
any allocation is attempted downstream (Section 3.6.1). 571
 10. Structured zone with plain FaceCenter through the fixed cg_sol_write (Sec- 572
tion 3.6.2): must still be accepted for CellDimension ≥ 3 (matching cg_- 573
gridlocation_write's existing behavior) and rejected below that, pinning 574
down the delegated-cgi_check_location behavior in both directions. 575

Note on cgnscheck regression risk. Once `get_data_size()` is fixed to honor an explicit point set (Section 2.5.2), files that previously produced a false-positive `ERROR:` for *any* reason unrelated to sizing will, for the first time, actually have their declared array size checked against `ListLength[]`. Any existing file whose `FlowSolution_t` array size does not genuinely match its point set’s `ListLength` — previously masked by the unconditional unstructured-zone bailout — will newly and correctly fail. This is the widest legacy-system ripple in this proposal: it is a correctness fix, not a regression, but should be called out explicitly in `cgnscheck` release notes so it is not mistaken for one.

9 Design Decisions

9.1 Alternative 1: Restrict the Relaxation to NGON_n Coverage Only

A narrower version of this proposal would restrict the relaxation to zones with `NGON_n` coverage, on the grounds that only `NGON_n` carries SIDS’s “all the faces in the grid” guarantee, and that treating classical 2-D sections’ combined `ElementSize` as “the face count of the zone” would misinterpret boundary-patch-only files. **Rejected.** This reasoning rests on reading “all faces” as “all geometric faces of the mesh” — but that is not, and never was, what face-centered data can address. An explicit `PointList` can only reference faces that carry element numbers, so on a classical mesh the reachable set is the section-defined faces regardless of whether the point set is present (Section 2.3). Once “complete” is read correctly as “complete over the faces the zone defines,” the coincidental-size-match objection dissolves: the rule in Section 3.2 never compares sizes against an external total at all; it requires the face sections’ `ElementRanges` to form one structurally verifiable contiguous block, which either holds or does not. The identity point set is exactly as redundant for a boundary `QUAD_4` block as for an `NGON_n` block, and the summation/contiguity machinery is identical — multiple `NGON_n` sections already required it (Section 3.2).

9.2 Alternative 2: New Explicit Marker Node (e.g., `FacesComplete`)

Introduce a boolean/flag child of `Zone_t` or `Elements_t` stating explicitly that a section enumerates all faces — reviving, in spirit, the 2005-era “`ZoneElementType`” idea from the original face-based storage proposal that predated `NGON_n/NFACE_n` standardization. **Rejected.** The element sections already carry the needed meaning by SIDS definition (§7.3); adding a redundant marker node reintroduces exactly the kind of “assert something already implied” overhead this CPEX exists to remove, just relocated from `FlowSolution_t` to `Zone_t`.

9.3 Alternative 3 (chosen): Structural Inference from Face-Section Coverage 610 611

Infer coverage structurally from the conditions in Section 3.2, using only guarantees 612
the SIDS already makes. No new nodes; no new trust assumptions beyond ones 613
the standard already relies on for `CellCenter`. This is the design adopted by this 614
CPEX. 615

9.4 Why `NGON_n` Takes Precedence in Mixed Zones (Tier Rule) 616

The one place where “sum all the face sections” would go wrong is the common 617
compatibility pattern of Example 5.5: an `NGON_n` block covering every face *plus* 618
classical boundary-patch sections duplicating the boundary faces for `BC_t/tool` 619
consumption. A naive union over all face sections would count those boundary 620
faces twice and silently change the implicit array’s length and meaning for precisely 621
this common polyhedral-plus-BC pattern. The two-tier rule in Section 3.2 resolves 622
this without any new markers: when `NGON_n` is present it is, by its SIDS definition, 623
already the complete face set, so any coexisting classical sections are necessarily 624
duplicates and are excluded from the indexing block; when no `NGON_n` exists, the 625
classical sections are the only face definitions there are, and their union is the block. 626
Each tier’s justification comes directly from existing SIDS text, and the two tiers 627
never overlap. 628

9.5 On Scope: Why Not Also Amend `ZoneSubRegion_t`? 629

`ZoneSubRegion_t` (§7.9) shares CPEX 0031 ancestry with `FlowSolution_t` and 630
`DiscreteData_t`, but it always requires an explicit region descriptor by design — a 631
subregion that happened to cover every face would be indistinguishable from simply 632
using `FlowSolution_t`, so there is no analogous “identity point set” pain point to 633
remove. This is left as an open question for the committee rather than folded into 634
this proposal’s scope. 635

10 Open Questions 636

1. Should `cgnscheck` *warn* (rather than silently accept) when it encounters the 637
pre-CPEX-0051 `identity-PointRange` pattern on a zone with complete face- 638
section coverage, to nudge writers toward the more compact implicit form? 639
This proposal takes no position; it is a tooling policy choice, not a conformance 640
question. 641
2. Should version gating use CPEX 0049’s `CGNSMinRequiredVersion_t`/feature- 642
mask mechanism, if adopted, or a plain `CGNSLibraryVersion_t` bump? This 643

- CPEX assumes whichever mechanism is current SIDS practice at the time of adoption and does not mandate a specific one.
3. CPEX 0050 (`DofStorage_t`), Example 4, already writes a `FaceCenter TraceSolution` via a single `cg_sol_write` call with no accompanying point set, for a zone implied to have `NGON_n` coverage. That example is only fully conformant once both this CPEX (for the missing point set) and the `cg_sol_write` allow-list fix (Section 2.5.1) are adopted. If both CPEXes proceed, their reference implementations should be sequenced so CPEX 0051 lands first, or the dependency should be noted explicitly in CPEX 0050.
 4. Should a companion CPEX add `IEdgeCenter/JEdgeCenter` to `GridLocation_t` for 2-D structured zones, restoring symmetry with 3-D's `IFaceCenter/JFaceCenter/KFaceCenter`? Raised during pre-submission review (Section 3.7); it requires a new enumeration value, which this CPEX deliberately avoids, so it is recorded here rather than folded in.

11 Summary of API Additions

The two implementation gaps from Section 2.5 (`cg_sol_write`'s allow-list, and `cgnscheck`'s missing point-set read) are pre-existing bugs, independent of the face-section-coverage logic itself, and do not need to wait on this CPEX's SIDS-level relaxation or on v5.0.0. They can and should land as standalone patches against the current release line as soon as they are reviewed; this CPEX's reference implementation only *depends on* them, so that Example 5.2's single-call write works end to end.

Symbol	Change	Release
<code>cg_sol_write</code>	Delegates location check to <code>cgi_check_location</code> , accepting <code>FaceCenter</code> directly (Sec. 3.6.2)	any (backportable)
<code>cg_discrete_write</code>	Same fix, mirrored	any (backportable)
<code>get_data_size()</code> (<code>cgnscheck</code>)	Point-set-aware sizing (Sec. 2.5.2)	any (backportable)
<code>cgi_zone_face_range()</code>	New internal helper (Sec. 3.6.1)	v5.0.0 (this CPEX)
<code>cg_zone_face_range()</code>	New public wrapper, read-side coverage check (Sec. 3.6.1)	v5.0.0 (this CPEX)

Aside from the new public `cg_zone_face_range()` wrapper, no new public enumerations, structs, or function signatures are introduced; `FaceCenter` already exists in

`GridLocation_t`. 668

12 References 669

1. CGNS SIDS, §7.3 `Elements_t` (element numbering, `NGON_n/NFACE_n` definition). 670
2. CGNS SIDS, §7.7 `FlowSolution_t` (Notes and `DataSize[]` function). 671
3. CGNS SIDS, §12.4 `DiscreteData_t`. 672
4. CGNS SIDS Changelog, Version 3.1 (CPEX 0031: `PointList/PointRange` revamp; `NGON_n/NFACE_n` formalization). 673
674
5. CPEX 0049, CGNS Compatibility Version Control (version-gating mechanism 675
referenced in Section 4). 676
6. CPEX 0050, DOF Storage Type (Example 4 dependency noted in Open Ques- 677
tion 3, Section 10). 678
7. S. Allmaras, “Face-Based Proposal” (2005-era precursor discussion of face-based 679
unstructured storage, predating standardized `NGON_n/NFACE_n`). 680